

# Student\_Prediction

February 10, 2026

```
[4]: """
Student Placement Prediction - Machine Learning Classification Model
Complete Pipeline: Data Analysis, Model Training, Evaluation & Visualization
"""

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score,
    GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, classification_report,
    roc_curve, auc, roc_auc_score)
import warnings
warnings.filterwarnings('ignore')

# Set style
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_palette("husl")

print("=="*90)
print("STUDENT PLACEMENT PREDICTION - MACHINE LEARNING CLASSIFICATION")
print("=="*90)

# =====
# 1. DATA LOADING AND EXPLORATION
# =====
print("\n" + "=="*90)
print("STEP 1: DATA LOADING AND EXPLORATION")
```

```

print("=="*90)

df = pd.read_csv('/home/nmit/Documents/Student_Placement.csv')

print(f"\n Dataset Shape: {df.shape}")
print(f" - Total Students: {df.shape[0]}")
print(f" - Total Features: {df.shape[1]}")

print("\n First 10 Rows:")
print(df.head(10))

print("\n Dataset Information:")
print(df.info())

print("\n Statistical Summary:")
print(df.describe())

print("\n Missing Values:")
print(df.isnull().sum())

print("\n Target Variable Distribution:")
print(df['Placement_Status'].value_counts())
print(f"\nPlacement Rate: {(df['Placement_Status'].str.contains('Placed', u
    ↴case=False).sum() / len(df) * 100):.2f}%")


# =====
# 2. DATA PREPROCESSING
# =====

print("\n" + "=="*90)
print("STEP 2: DATA PREPROCESSING")
print("=="*90)

# Handle the target variable (fix inconsistency in data)
df['Placement_Status'] = df['Placement_Status'].str.strip()
df['Placement_Status'] = df['Placement_Status'].replace('Not Place', 'NotU
    ↴Placed')

print("\n Cleaned Placement Status")
print(f"Unique values: {df['Placement_Status'].unique()}")

# Encode target variable
le = LabelEncoder()
df['Placement_Encoded'] = le.fit_transform(df['Placement_Status'])
print(f"\n Label Encoding: {dict(zip(le.classes_, le.transform(le.
    ↴classes_)))}")


# Separate features and target

```

```

X = df.drop(['Placement_Status', 'Placement_Encoded'], axis=1)
y = df['Placement_Encoded']

print(f"\n Feature Matrix Shape: {X.shape}")
print(f" Target Vector Shape: {y.shape}")
print(f"\nFeature Names: {list(X.columns)}")

# =====
# 3. TRAIN-TEST SPLIT
# =====
print("\n" + "="*90)
print("STEP 3: TRAIN-TEST SPLIT")
print("="*90)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

print(f"\n Training Set Size: {X_train.shape[0]} samples ({(len(X_train)/
    len(X)*100):.0f}%)")
print(f" Test Set Size: {X_test.shape[0]} samples ({(len(X_test)/len(X)*100):.
    0f}%)")
print(f"\nTraining set class distribution:")
print(pd.Series(y_train).value_counts())
print(f"\nTest set class distribution:")
print(pd.Series(y_test).value_counts())

# =====
# 4. FEATURE SCALING
# =====
print("\n" + "="*90)
print("STEP 4: FEATURE SCALING")
print("="*90)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("\n Features Standardized (Mean=0, Std=1)")
print("\nScaling Parameters:")
for i, col in enumerate(X.columns):
    print(f" {col}: Mean={scaler.mean_[i]:.2f}, Std={scaler.scale_[i]:.2f}")

# =====
# 5. MODEL TRAINING - MULTIPLE ALGORITHMS
# =====
print("\n" + "="*90)

```

```

print("STEP 5: TRAINING MULTIPLE CLASSIFICATION MODELS")
print("=="*90)

# Define models
models = {
    'Logistic Regression': LogisticRegression(random_state=42, max_iter=1000),
    'Decision Tree': DecisionTreeClassifier(random_state=42, max_depth=5),
    'Random Forest': RandomForestClassifier(n_estimators=100, random_state=42),
    'Support Vector Machine': SVC(kernel='rbf', random_state=42,
        probability=True),
    'K-Nearest Neighbors': KNeighborsClassifier(n_neighbors=3),
    'Naive Bayes': GaussianNB(),
    'Gradient Boosting': GradientBoostingClassifier(n_estimators=100,
        random_state=42)
}

# Storage for results
results = []
trained_models = {}

print("\n Training Models...")
for name, model in models.items():
    print(f"\n Training {name}...", end=" ")

    # Train model
    model.fit(X_train_scaled, y_train)
    trained_models[name] = model

    # Make predictions
    y_pred_train = model.predict(X_train_scaled)
    y_pred_test = model.predict(X_test_scaled)

    # Calculate metrics
    train_accuracy = accuracy_score(y_train, y_pred_train)
    test_accuracy = accuracy_score(y_test, y_pred_test)
    precision = precision_score(y_test, y_pred_test, zero_division=0)
    recall = recall_score(y_test, y_pred_test, zero_division=0)
    f1 = f1_score(y_test, y_pred_test, zero_division=0)

    # Cross-validation score
    cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=3,
        scoring='accuracy')
    cv_mean = cv_scores.mean()

    results.append({
        'Model': name,
        'Train Accuracy': train_accuracy,

```

```

        'Test Accuracy': test_accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-Score': f1,
        'CV Score': cv_mean
    })

    print(f" (Accuracy: {test_accuracy:.4f}))"

# Convert results to DataFrame
results_df = pd.DataFrame(results)
results_df = results_df.sort_values('Test Accuracy', ascending=False) .
    ↪reset_index(drop=True)

print("\n" + "="*90)
print("MODEL PERFORMANCE COMPARISON")
print("="*90)
print(results_df.to_string(index=False))

# Select best model
best_model_name = results_df.iloc[0]['Model']
best_model = trained_models[best_model_name]
print(f"\n BEST MODEL: {best_model_name}")
print(f" Test Accuracy: {results_df.iloc[0]['Test Accuracy']:.4f}")

# =====
# 6. DETAILED EVALUATION OF BEST MODEL
# =====

print("\n" + "="*90)
print(f"STEP 6: DETAILED EVALUATION - {best_model_name}")
print("="*90)

# Predictions
y_pred = best_model.predict(X_test_scaled)
y_pred_proba = best_model.predict_proba(X_test_scaled)[:, 1] if ↪
    hasattr(best_model, 'predict_proba') else None

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("\n Confusion Matrix:")
print(cm)
print(f"\nTrue Negatives (TN): {cm[0,0]}")
print(f"False Positives (FP): {cm[0,1]}")
print(f"False Negatives (FN): {cm[1,0]}")
print(f"True Positives (TP): {cm[1,1]}")

# Detailed metrics

```

```

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, zero_division=0)
recall = recall_score(y_test, y_pred, zero_division=0)
f1 = f1_score(y_test, y_pred, zero_division=0)

print("\n Performance Metrics:")
print(f"    Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f"    Precision: {precision:.4f} ({precision*100:.2f}%)")
print(f"    Recall: {recall:.4f} ({recall*100:.2f}%)")
print(f"    F1-Score: {f1:.4f} ({f1*100:.2f}%)")

if y_pred_proba is not None:
    roc_auc = roc_auc_score(y_test, y_pred_proba)
    print(f"    ROC-AUC: {roc_auc:.4f} ({roc_auc*100:.2f}%)")

print("\n Classification Report:")
print(classification_report(y_test, y_pred,
                           target_names=['Not Placed', 'Placed'],
                           zero_division=0))

# Feature Importance (if available)
if hasattr(best_model, 'feature_importances_'):
    print("\n Feature Importance:")
    feature_importance = pd.DataFrame({
        'Feature': X.columns,
        'Importance': best_model.feature_importances_
    }).sort_values('Importance', ascending=False)
    print(feature_importance.to_string(index=False))
elif hasattr(best_model, 'coef_'):
    print("\n Feature Coefficients:")
    feature_coef = pd.DataFrame({
        'Feature': X.columns,
        'Coefficient': best_model.coef_[0]
    }).sort_values('Coefficient', ascending=False)
    print(feature_coef.to_string(index=False))

# =====
# 7. VISUALIZATIONS
# =====

print("\n" + "="*90)
print("STEP 7: GENERATING VISUALIZATIONS")
print("="*90)

# Create comprehensive visualization
fig = plt.figure(figsize=(20, 16))

# 1. Model Comparison - Accuracy

```

```

ax1 = plt.subplot(3, 3, 1)
models_sorted = results_df.sort_values('Test Accuracy', ascending=True)
colors = ['#2ecc71' if x == best_model_name else '#3498db' for x in
         models_sorted['Model']]
bars = plt.barh(models_sorted['Model'], models_sorted['Test Accuracy'], color=colors, edgecolor='black', linewidth=1.5)
plt.xlabel('Accuracy', fontsize=12, fontweight='bold')
plt.title('Model Accuracy Comparison', fontsize=14, fontweight='bold')
plt.xlim([0, 1.1])
for i, bar in enumerate(bars):
    width = bar.get_width()
    plt.text(width + 0.02, bar.get_y() + bar.get_height()/2,
              f'{width:.3f}', ha='left', va='center', fontweight='bold')
plt.axvline(x=0.8, color='red', linestyle='--', alpha=0.5, label='80% Threshold')
plt.legend()

# 2. Confusion Matrix Heatmap
ax2 = plt.subplot(3, 3, 2)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', square=True,
            xticklabels=['Not Placed', 'Placed'],
            yticklabels=['Not Placed', 'Placed'],
            cbar_kws={'label': 'Count'}, linewidths=2, linecolor='black')
plt.ylabel('Actual', fontsize=12, fontweight='bold')
plt.xlabel('Predicted', fontsize=12, fontweight='bold')
plt.title(f'Confusion Matrix - {best_model_name}', fontsize=14,
          fontweight='bold')

# 3. Performance Metrics Bar Chart
ax3 = plt.subplot(3, 3, 3)
metrics_data = {
    'Accuracy': accuracy,
    'Precision': precision,
    'Recall': recall,
    'F1-Score': f1
}
bars = plt.bar(metrics_data.keys(), metrics_data.values(),
               color=['#e74c3c', '#3498db', '#2ecc71', '#f39c12'],
               edgecolor='black', linewidth=2, alpha=0.8)
plt.ylabel('Score', fontsize=12, fontweight='bold')
plt.title(f'Performance Metrics - {best_model_name}', fontsize=14,
          fontweight='bold')
plt.ylim([0, 1.1])
plt.axhline(y=0.8, color='red', linestyle='--', alpha=0.5)
for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2., height + 0.02,

```

```

        f'{height:.3f}', ha='center', va='bottom', fontweight='bold', u
        fontsize=11)

# 4. ROC Curve (if available)
ax4 = plt.subplot(3, 3, 4)
if y_pred_proba is not None:
    fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, color='#e74c3c', lw=3, label=f'{best_model_name} (AUC = u
    {roc_auc:.3f})')
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random u
    Classifier')
    plt.fill_between(fpr, tpr, alpha=0.2, color="#e74c3c")
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate', fontsize=12, fontweight='bold')
    plt.ylabel('True Positive Rate', fontsize=12, fontweight='bold')
    plt.title('ROC Curve', fontsize=14, fontweight='bold')
    plt.legend(loc="lower right", fontsize=10)
    plt.grid(True, alpha=0.3)
else:
    plt.text(0.5, 0.5, 'ROC Curve\nNot Available\nfor this model',
             ha='center', va='center', fontsize=14)
    plt.axis('off')

# 5. Train vs Test Accuracy
ax5 = plt.subplot(3, 3, 5)
x = np.arange(len(results_df))
width = 0.35
bars1 = plt.bar(x - width/2, results_df['Train Accuracy'], width,
                 label='Train', color='#3498db', edgecolor='black', linewidth=1.5)
bars2 = plt.bar(x + width/2, results_df['Test Accuracy'], width,
                 label='Test', color='#2ecc71', edgecolor='black', linewidth=1.5)
plt.xlabel('Models', fontsize=12, fontweight='bold')
plt.ylabel('Accuracy', fontsize=12, fontweight='bold')
plt.title('Train vs Test Accuracy (Overfitting Check)', fontsize=14, u
        fontweight='bold')
plt.xticks(x, results_df['Model'], rotation=45, ha='right')
plt.legend()
plt.grid(True, alpha=0.3, axis='y')
plt.tight_layout()

# 6. Feature Importance (if available)
ax6 = plt.subplot(3, 3, 6)
if hasattr(best_model, 'feature_importances_'):
    feature_imp = pd.DataFrame({
        'Feature': X.columns,

```

```

        'Importance': best_model.feature_importances_
    }).sort_values('Importance', ascending=True)
colors_imp = plt.cm.viridis(np.linspace(0, 1, len(feature_imp)))
bars = plt.barih(feature_imp['Feature'], feature_imp['Importance'],
                  color=colors_imp, edgecolor='black', linewidth=1.5)
plt.xlabel('Importance', fontsize=12, fontweight='bold')
plt.title('Feature Importance', fontsize=14, fontweight='bold')
for i, bar in enumerate(bars):
    width = bar.get_width()
    plt.text(width + 0.01, bar.get_y() + bar.get_height()/2,
              f'{width:.3f}', ha='left', va='center', fontweight='bold')
elif hasattr(best_model, 'coef_'):
    feature_coef = pd.DataFrame({
        'Feature': X.columns,
        'Coefficient': np.abs(best_model.coef_[0])
    }).sort_values('Coefficient', ascending=True)
    colors_coef = plt.cm.plasma(np.linspace(0, 1, len(feature_coef)))
    bars = plt.barih(feature_coef['Feature'], feature_coef['Coefficient'],
                      color=colors_coef, edgecolor='black', linewidth=1.5)
    plt.xlabel('Absolute Coefficient Value', fontsize=12, fontweight='bold')
    plt.title('Feature Coefficients', fontsize=14, fontweight='bold')
    for i, bar in enumerate(bars):
        width = bar.get_width()
        plt.text(width + 0.01, bar.get_y() + bar.get_height()/2,
                  f'{width:.3f}', ha='left', va='center', fontweight='bold')
else:
    plt.text(0.5, 0.5, 'Feature Importance\nNot Available\nfor this model',
            ha='center', va='center', fontsize=14)
    plt.axis('off')

# 7. Cross-Validation Scores
ax7 = plt.subplot(3, 3, 7)
models_cv = results_df.sort_values('CV Score', ascending=True)
colors_cv = ['#2ecc71' if x == best_model_name else '#e67e22' for x in
            models_cv['Model']]
bars = plt.barih(models_cv['Model'], models_cv['CV Score'], color=colors_cv,
                  edgecolor='black', linewidth=1.5)
plt.xlabel('Cross-Validation Score', fontsize=12, fontweight='bold')
plt.title('3-Fold Cross-Validation Scores', fontsize=14, fontweight='bold')
plt.xlim([0, 1.1])
for i, bar in enumerate(bars):
    width = bar.get_width()
    plt.text(width + 0.02, bar.get_y() + bar.get_height()/2,
              f'{width:.3f}', ha='left', va='center', fontweight='bold')

# 8. Target Distribution
ax8 = plt.subplot(3, 3, 8)

```

```

placement_counts = df['Placement_Status'].value_counts()
colors_pie = ['#e74c3c', '#2ecc71']
wedges, texts, autotexts = plt.pie(placement_counts.values,
    labels=placement_counts.index,
    autopct='%.1f%%', startangle=90,
    colors=colors_pie,
    wedgeprops={'edgecolor': 'black',
    'linewidth': 2},
    textprops={'fontsize': 12, 'fontweight': 'bold'})
plt.title('Target Variable Distribution', fontsize=14, fontweight='bold')

# 9. Correlation Heatmap
ax9 = plt.subplot(3, 3, 9)
correlation_matrix = df[list(X.columns) + ['Placement_Encoded']].corr()
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm',
    center=0, square=True, linewidths=1, linecolor='black',
    cbar_kws={'label': 'Correlation'})
plt.title('Feature Correlation Matrix', fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)

plt.suptitle('Student Placement Prediction - Comprehensive Analysis',
    fontsize=18, fontweight='bold', y=0.995)
plt.tight_layout()
plt.savefig('/home/nmit/Pictures/placement_prediction_comprehensive.png',
    dpi=300, bbox_inches='tight')
print("\n Comprehensive visualization saved!")

# =====
# Additional Visualization: Decision Tree (if best model is tree-based)
# =====
if 'Decision Tree' in best_model_name or 'Random Forest' in best_model_name:
    fig2 = plt.figure(figsize=(20, 12))

    if 'Decision Tree' in best_model_name:
        tree_model = best_model
    else:
        tree_model = best_model.estimators_[0]

    plot_tree(tree_model, feature_names=X.columns,
        class_names=['Not Placed', 'Placed'],
        filled=True, rounded=True, fontsize=10)
    plt.title(f'Decision Tree Visualization - {best_model_name}',
        fontsize=16, fontweight='bold', pad=20)
    plt.savefig('/home/nmit/Pictures/decision_tree_visualization.png',
        dpi=300, bbox_inches='tight')

```

```

    print(" Decision tree visualization saved!")

# =====
# Prediction Examples
# =====

print("\n" + "="*90)
print("STEP 8: SAMPLE PREDICTIONS")
print("="*90)

print("\n Sample Predictions on Test Data:")
sample_predictions = pd.DataFrame({
    'Actual': ['Placed' if y == 1 else 'Not Placed' for y in y_test[:5]],
    'Predicted': ['Placed' if y == 1 else 'Not Placed' for y in y_pred[:5]],
    'Correct': [ ' ' if a == p else ' ' for a, p in zip(y_test[:5], y_pred[:5])]
})
print(sample_predictions.to_string(index=False))

# =====
# Save Results
# =====

print("\n" + "="*90)
print("STEP 9: SAVING RESULTS")
print("="*90)

# Save model comparison
results_df.to_csv('/home/nmit/Pictures/model_comparison_results.csv',□
    ↪index=False)
print(" Model comparison results saved!")

# Save detailed metrics
with open('/home/nmit/Pictures/placement_prediction_report.txt', 'w') as f:
    f.write(" "*90 + "\n")
    f.write("STUDENT PLACEMENT PREDICTION - DETAILED REPORT\n")
    f.write(" "*90 + "\n\n")

    f.write("DATASET SUMMARY\n")
    f.write("-"*90 + "\n")
    f.write(f"Total Students: {df.shape[0]}\n")
    f.write(f"Features: {', '.join(X.columns)}\n")
    f.write(f"Placement Rate: {(df['Placement_Status'].str.contains('Placed',□
        ↪case=False).sum() / len(df) * 100):.2f}%\n\n")

    f.write("MODEL COMPARISON\n")
    f.write("-"*90 + "\n")
    f.write(results_df.to_string(index=False))
    f.write("\n\n")

```

```

f.write(f"BEST MODEL: {best_model_name}\n")
f.write("-"*90 + "\n")
f.write(f"Test Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%) \n")
f.write(f"Precision: {precision:.4f} ({precision*100:.2f}%) \n")
f.write(f"Recall: {recall:.4f} ({recall*100:.2f}%) \n")
f.write(f"F1-Score: {f1:.4f} ({f1*100:.2f}%) \n\n")

f.write("CONFUSION MATRIX\n")
f.write("-"*90 + "\n")
f.write(f"True Negatives (TN): {cm[0,0]}\n")
f.write(f"False Positives (FP): {cm[0,1]}\n")
f.write(f"False Negatives (FN): {cm[1,0]}\n")
f.write(f"True Positives (TP): {cm[1,1]}\n\n")

f.write("CLASSIFICATION REPORT\n")
f.write("-"*90 + "\n")
f.write(classification_report(y_test, y_pred,
                             target_names=['Not Placed', 'Placed'],
                             zero_division=0))
f.write("\n")

if hasattr(best_model, 'feature_importances_'):
    f.write("FEATURE IMPORTANCE\n")
    f.write("-"*90 + "\n")
    f.write(feature_importance.to_string(index=False))
    f.write("\n")

print(" Detailed report saved!")

# Save predictions
predictions_df = pd.DataFrame({
    'Actual': y_test,
    'Predicted': y_pred,
    'Actual_Label': ['Placed' if y == 1 else 'Not Placed' for y in y_test],
    'Predicted_Label': ['Placed' if y == 1 else 'Not Placed' for y in y_pred]
})
predictions_df.to_csv('/home/nmit/Pictures/test_predictions.csv', index=False)
print(" Test predictions saved!")

print("\n" + "="*90)
print("ANALYSIS COMPLETE!")
print("=*90)
print("\n Generated Files:")
print(" 1. placement_prediction_comprehensive.png - Main visualizations")
print(" 2. model_comparison_results.csv - Model performance comparison")
print(" 3. placement_prediction_report.txt - Detailed analysis report")
print(" 4. test_predictions.csv - Predictions on test data")

```

```
if 'Decision Tree' in best_model_name or 'Random Forest' in best_model_name:  
    print("  5. decision_tree_visualization.png - Decision tree structure")  
  
print("\n" + "="*90)  
print("  KEY FINDINGS")  
print("="*90)  
print(f"  Best Model: {best_model_name}")  
print(f"  Test Accuracy: {accuracy*100:.2f}%")  
print(f"  Precision: {precision*100:.2f}%")  
print(f"  Recall: {recall*100:.2f}%")  
print(f"  F1-Score: {f1*100:.2f}%")  
print("\nThe model can predict student placement with high accuracy!")  
print("="*90)
```

STUDENT PLACEMENT PREDICTION = MACHINE LEARNING CLASSIFICATION

## STEP 1: DATA LOADING AND EXPLORATION

Dataset Shape: (10, 6)  
- Total Students: 10  
- Total Features: 6

First 10 Rows:

	CGPA	Internships	Technical_Skills	Communication_Skills	Aptitude_Score	\
0	6.5	0	60	55	58	
1	7.2	1	70	65	68	
2	8.1	2	85	80	82	
3	8.8	3	90	85	90	
4	5.9	0	55	50	55	
5	7.5	1	75	70	72	
6	9.0	3	95	90	92	
7	6.8	1	65	60	62	
8	8.5	2	88	82	85	
9	6.0	0	58	52	57	

	Placement_Status
0	Not Place
1	Placed
2	Placed
3	Placed

```

4      Not Placed
5      Placed
6      Placed
7      Not Placed
8      Placed
9      Not Placed

Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CGPA             10 non-null    float64
 1   Internships       10 non-null    int64  
 2   Technical_Skills 10 non-null    int64  
 3   Communication_Skills 10 non-null  int64  
 4   Aptitutde_Score   10 non-null    int64  
 5   Placement_Status  10 non-null    object 
dtypes: float64(1), int64(4), object(1)
memory usage: 612.0+ bytes
None

Statistical Summary:
          CGPA   Internships  Technical_Skills  Communication_Skills \
count  10.000000  10.000000  10.000000  10.000000
mean   7.430000  1.300000  74.100000  68.900000
std    1.137297  1.159502  14.624561  14.647336
min   5.900000  0.000000  55.000000  50.000000
25%   6.575000  0.250000  61.250000  56.250000
50%   7.350000  1.000000  72.500000  67.500000
75%   8.400000  2.000000  87.250000  81.500000
max   9.000000  3.000000  95.000000  90.000000

          Aptitutde_Score
count      10.000000
mean      72.100000
std       14.216188
min      55.000000
25%     59.000000
50%     70.000000
75%     84.250000
max     92.000000

Missing Values:
CGPA          0
Internships    0
Technical_Skills 0

```

```
Communication_Skills      0
Aptitutde_Score          0
Placement_Status          0
dtype: int64

    Target Variable Distribution:
Placement_Status
Placed      6
Not Placed   3
Not Place    1
Name: count, dtype: int64

Placement Rate: 90.00%
=====
=====
STEP 2: DATA PREPROCESSING
=====
=====

    Cleaned Placement Status
Unique values: ['Not Placed' 'Placed']

Label Encoding: {'Not Placed': 0, 'Placed': 1}

Feature Matrix Shape: (10, 5)
Target Vector Shape: (10,)

Feature Names: ['CGPA', 'Internships', 'Technical_Skills',
'Communication_Skills', 'Aptitutde_Score']
=====
=====
STEP 3: TRAIN-TEST SPLIT
=====
=====

Training Set Size: 7 samples (70%)
Test Set Size: 3 samples (30%)

Training set class distribution:
Placement_Encoded
1      4
0      3
Name: count, dtype: int64

Test set class distribution:
Placement_Encoded
```

```
1    2  
0    1  
Name: count, dtype: int64
```

```
=====  
=====  
STEP 4: FEATURE SCALING  
=====  
=====
```

```
Features Standardized (Mean=0, Std=1)
```

```
Scaling Parameters:
```

```
CGPA: Mean=7.10, Std=0.94  
Internships: Mean=0.86, Std=0.83  
Technical_Skills: Mean=70.14, Std=12.21  
Communication_Skills: Mean=64.86, Std=12.15  
Aptitudine_Score: Mean=68.14, Std=11.28
```

```
=====  
=====  
STEP 5: TRAINING MULTIPLE CLASSIFICATION MODELS  
=====  
=====
```

```
Training Models...
```

```
Training Logistic Regression... (Accuracy: 0.6667)
```

```
Training Decision Tree... (Accuracy: 0.6667)
```

```
(Accuracy: 1.0000)rest...
```

```
Training Support Vector Machine... (Accuracy: 0.6667)
```

```
Training K-Nearest Neighbors... (Accuracy: 0.6667)
```

```
Training Naive Bayes... (Accuracy: 0.6667)
```

```
Training Gradient Boosting... (Accuracy: 0.6667)
```

```
=====  
=====  
MODEL PERFORMANCE COMPARISON  
=====  
=====
```

	Model	Train Accuracy	Test Accuracy	Precision	Recall
F1-Score	CV Score				

	Random Forest	1.0	1.000000	1.000000	1.0
1.0	1.000000				
	Logistic Regression	1.0	0.666667	0.666667	1.0
0.8	1.000000				
	Decision Tree	1.0	0.666667	0.666667	1.0
0.8	1.000000				
	Support Vector Machine	1.0	0.666667	0.666667	1.0
0.8	1.000000				
	K-Nearest Neighbors	1.0	0.666667	0.666667	1.0
0.8	1.000000				
	Naive Bayes	1.0	0.666667	0.666667	1.0
0.8	0.833333				
	Gradient Boosting	1.0	0.666667	0.666667	1.0
0.8	1.000000				

BEST MODEL: Random Forest

Test Accuracy: 1.0000

---



---



---

STEP 6: DETAILED EVALUATION - Random Forest

---



---



---

Confusion Matrix:

```
[[1 0]
 [0 2]]
```

True Negatives (TN): 1  
 False Positives (FP): 0  
 False Negatives (FN): 0  
 True Positives (TP): 2

Performance Metrics:

Accuracy: 1.0000 (100.00%)  
 Precision: 1.0000 (100.00%)  
 Recall: 1.0000 (100.00%)  
 F1-Score: 1.0000 (100.00%)  
 ROC-AUC: 1.0000 (100.00%)

Classification Report:

	precision	recall	f1-score	support
Not Placed	1.00	1.00	1.00	1
Placed	1.00	1.00	1.00	2
accuracy			1.00	3
macro avg	1.00	1.00	1.00	3

weighted avg 1.00 1.00 1.00 3

Feature Importance:

Feature	Importance
CGPA	0.224490
Communication_Skills	0.224490
Aptitudde_Score	0.204082
Internships	0.193878
Technical_Skills	0.153061

=====

=====

=====

STEP 7: GENERATING VISUALIZATIONS

=====

=====

Comprehensive visualization saved!  
Decision tree visualization saved!

=====

=====

=====

STEP 8: SAMPLE PREDICTIONS

=====

=====

Sample Predictions on Test Data:

Actual	Predicted	Correct
Not Placed	Not Placed	
Placed	Placed	
Placed	Placed	

=====

=====

=====

STEP 9: SAVING RESULTS

=====

=====

Model comparison results saved!  
Detailed report saved!  
Test predictions saved!

=====

=====

=====

ANALYSIS COMPLETE!

=====

=====

Generated Files:

1. placement\_prediction\_comprehensive.png - Main visualizations
  2. model\_comparison\_results.csv - Model performance comparison
  3. placement\_prediction\_report.txt - Detailed analysis report
  4. test\_predictions.csv - Predictions on test data
  5. decision\_tree\_visualization.png - Decision tree structure
- 
- 
- 

## KEY FINDINGS

---



---

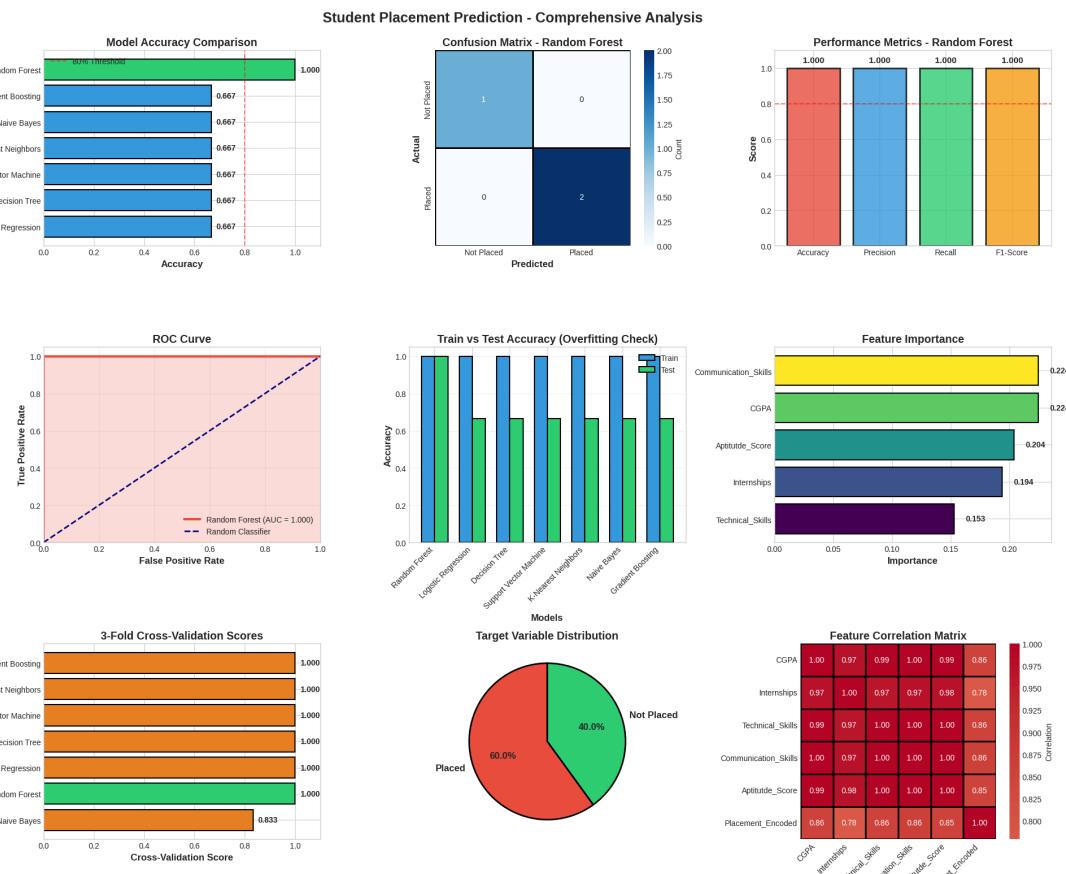
Best Model: Random Forest  
 Test Accuracy: 100.00%  
 Precision: 100.00%  
 Recall: 100.00%  
 F1-Score: 100.00%

The model can predict student placement with high accuracy!

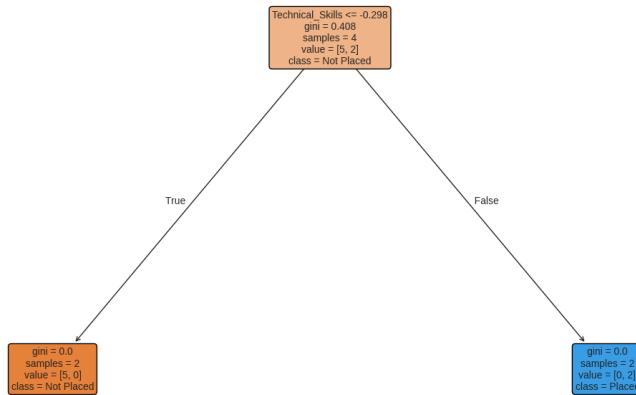
---



---



### Decision Tree Visualization - Random Forest



[ ]: